

Practical 8

Aim: - Building a Recurrent Neural Network (RNN) for Predicting Patient Readmission

Theory: -


Recurrent neural networks (RNNs) are a class of artificial neural network commonly used for sequential data processing. Unlike feedforward neural networks, which process data in a single pass, RNNs process data across multiple time steps, making them well-adapted for modelling and processing text, speech, and time series.^[1]

The building block of RNNs is the *recurrent unit*. This unit maintains a hidden state, essentially a form of memory, which is updated at each time step based on the current input and the previous hidden state. This feedback loop allows the network to learn from past inputs, and incorporate that knowledge into its current processing.

Early RNNs suffered from the vanishing gradient problem, limiting their ability to learn long-range dependencies. This was solved by the long short-term memory (LSTM) variant in 1997, thus making it the standard architecture for RNN.

RNNs have been applied to tasks such as unsegmented, connected handwriting recognition, speech recognition, natural language processing, and neural machine translation.

Block Diagram: -



Dataset Description: -

The **Hospital Readmission Dataset** analyzes patient demographics, medical history, and treatment details to predict hospital readmissions. Key attributes include age, time in hospital, number of lab procedures, medications, outpatient/inpatient/emergency visits, medical specialty, diagnoses, and test results (e.g., glucose and A1C tests). The target variable indicates whether a patient was readmitted within a specific timeframe (Yes or No). This dataset supports efforts to identify factors contributing to readmissions and improve healthcare outcomes.

Source code and Output: -

```
In [1]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler, OneHotEncoder
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, SimpleRNN, Dropout
from tensorflow.keras.utils import to_categorical
```

```
In [2]: data = pd.read_csv("C:/Users/DELL/Downloads/hospital_readmissions.csv")
print(data.head())
```

	age	time_in_hospital	n_lab_procedures	n_procedures	n_medications	\
0	[70-80)	8	72	1	18	
1	[70-80)	3	34	2	13	
2	[50-60)	5	45	0	18	
3	[70-80)	2	36	0	12	
4	[60-70)	1	42	0	7	

	n_outpatient	n_inpatient	n_emergency	medical_specialty	diag_1	\
0	2	0	0	Missing	Circulatory	
1	0	0	0	Other	Other	
2	0	0	0	Missing	Circulatory	
3	1	0	0	Missing	Circulatory	
4	0	0	0	InternalMedicine	Other	

	diag_2	diag_3	glucose_test	A1Ctest	change	diabetes_med	\
0	Respiratory	Other	no	no	no	yes	

```
In [3]: age_mapping = {
    '0-10': 5,
    '10-20': 15,
    '20-30': 25,
    '30-40': 35,
    '40-50': 45,
    '50-60': 55,
    '60-70': 65,
    '70-80': 75,
    '80-90': 85,
    '90-100': 95
}

# Apply mapping to the 'age' column
data['age'] = data['age'].map(age_mapping)
```

```
In [4]: print(data.columns)

Index(['age', 'time_in_hospital', 'n_lab_procedures', 'n_procedures',
       'n_medications', 'n_outpatient', 'n_inpatient', 'n_emergency',
       'medical_specialty', 'diag_1', 'diag_2', 'diag_3', 'glucose_test',
       'A1Ctest', 'change', 'diabetes_med', 'readmitted'],
      dtype='object')
```

```
In [5]: categorical_columns = ['medical_specialty', 'diag_1', 'diag_2', 'diag_3', 'glucose_test',
                              'A1Ctest', 'change', 'diabetes_med']
data = pd.get_dummies(data, columns=categorical_columns, drop_first=True)
```

```
In [7]: features = data[['age', 'time_in_hospital', 'n_lab_procedures', 'n_medications', 'n_procedures',
                        'n_outpatient', 'n_inpatient', 'n_emergency'] + list(data.columns[data.columns.str.startswith(('medical_special',
                        'A1Ctest', 'change', 'diabetes_med'))])]

```

```
In [8]: target = data['readmitted']
```

```
In [9]: scaler = MinMaxScaler()
features_scaled = scaler.fit_transform(features)
```

```
In [15]: X = features_scaled.reshape(features_scaled.shape[0], 1, features_scaled.shape[1])
```

```
In [17]: y = data['readmitted'].map({'yes': 1, 'no': 0})

# Ensure no NaN values exist after mapping
if y.isnull().sum() > 0:
    raise ValueError("Target contains unmapped values or missing data!")
```

```

scaler = MinMaxScaler()
features_scaled = scaler.fit_transform(features)

# Reshape features for RNN input
X = features_scaled.reshape(features_scaled.shape[0], 1, features_scaled.shape[1])

# Verify target values
print("Unique values in y:", np.unique(y))

Unique values in y: [0 1]

```

```
In [18]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [19]: model = Sequential()

# Add RNN Layer
model.add(SimpleRNN(64, activation='relu', input_shape=(X_train.shape[1], X_train.shape[2])))

# Add Dropout Layer
model.add(Dropout(0.2))

# Add Dense output Layer
model.add(Dense(1, activation='sigmoid'))

# Compile the model
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

```

```
In [25]: history = model.fit(X_train, y_train, epochs=100, batch_size=32, validation_data=(X_test, y_test), verbose=1)
```

```

0.6050
Epoch 95/100
625/625 [=====] - 1s 2ms/step - loss: 0.6153 - accuracy: 0.6576 - val_loss: 0.6659 - val_accuracy:
0.6072
Epoch 96/100
625/625 [=====] - 1s 2ms/step - loss: 0.6160 - accuracy: 0.6525 - val_loss: 0.6672 - val_accuracy:
0.6062
Epoch 97/100
625/625 [=====] - 1s 2ms/step - loss: 0.6145 - accuracy: 0.6583 - val_loss: 0.6685 - val_accuracy:
0.6016
Epoch 98/100
625/625 [=====] - 1s 2ms/step - loss: 0.6163 - accuracy: 0.6530 - val_loss: 0.6658 - val_accuracy:
0.6074
Epoch 99/100
625/625 [=====] - 1s 2ms/step - loss: 0.6150 - accuracy: 0.6545 - val_loss: 0.6671 - val_accuracy:
0.6054
Epoch 100/100
625/625 [=====] - 1s 2ms/step - loss: 0.6137 - accuracy: 0.6556 - val_loss: 0.6666 - val_accuracy:
0.6088

```

```
In [27]: loss, accuracy = model.evaluate(X_test, y_test)
print(f"Test Accuracy: {accuracy * 100:.2f}%")
```

```
157/157 [=====] - 0s 907us/step - loss: 0.6666 - accuracy: 0.6088
```

```
In [28]: # Make predictions
predictions = (model.predict(X_test) > 0.5).astype("int32")
print("Predictions for the first 10 samples:", predictions[:10])
```

```

157/157 [=====] - 0s 838us/step
Predictions for the first 10 samples: [[1]
[0]
[0]
[0]
[1]
[0]
[1]
[0]
[0]]

```

Conclusion: - The practical demonstrates the potential of RNNs in predicting hospital readmissions by capturing sequential patterns in patient data, enhancing predictive accuracy for healthcare management.